

Agents

Gunhee Kim




서울대학교
SEOUL NATIONAL UNIVERSITY

Contents

- Agent Intro
- ChatGPT Work

An Era of AI Agents


- AI agents
 - Engage with their environments through continuous learning, reasoning, and adaptation
 - Autonomously executing complex tasks that once required human expertise
- Key components
 - LLMs (or VLMs): general-purpose processors, unifying perception, decision-making, and action
 - Memory: longitudinal experience accumulation
 - Tools: execute tasks and environmental interaction



OpenAI's Deep Research (2/5/2025)

- An AI agent powered by the o3 model
 - Autonomously find, analyze, and synthesize information from online sources
 - Built for intensive knowledge work in areas like finance, science, policy, and engineering and need thorough, precise, and reliable research
 - Trained on real-world tasks requiring browser and Python tool use
- How to use
 - In ChatGPT, enter a query with selecting 'deep research' in the message composer
 - 5 to 30 minutes to completion

What can I help with?



OpenAI's Deep Research Benchmarks

- Humanity's Last Exam

- Expert-level multiple-choice and short-answer questions across over 100 subjects, from linguistics and rocket science to ecology and mathematics

* Model is not multi-modal,
evaluated on text-only subset
**with browsing + python tools

- GAIA

- Real-world questions, requiring a combination of reasoning, web browsing, multimodal fluency, and tool-use proficiency

Model	Accuracy (%)
GPT-4o	3.3
Grok-2	3.8
Claude 3.5 Sonnet	4.3
Gemini Thinking	6.2
OpenAI o1	9.1
DeepSeek-R1*	9.4
o3-mini (medium)*	10.5
o3-mini (high)*	13.0
OpenAI deep research**	26.6

	Level 1	Level 2	Level 3	Avg.
Previous SOTA	67.92	67.44	42.31	63.64
Deep Research (pass@1)	74.29	69.06	47.6	67.36
Deep Research (cons@64)	78.66	73.21	58.03	72.57

<https://lastexam.ai/>


G. Mialon et al., GAIA: a benchmark for General AI Assistants, 2024

Harness Engineering

- Agent = model + harness
 - Agent harness: the software around an LLM: tools, memory, state, execution, guardrails, and observability
- How we build systems around models to turn them into work engines
- Behavior we want → Harness design to let the model achieve this


What Is an Agent Harness?

An agent combines a language model with a harness that surrounds and supports it.




Harness Engineering

- From Weights to Context to Harness
 - From parametric knowledge and prompting toward harness-level infrastructure such as tool ecosystems, protocols, skills, and multi-agent orchestration




Workflow

- A multi-step, dynamic process within a dynamic control graph
 - Orchestrate multiple API calls, tasks, agents, and even human-in-the-loop steps
 - Rather than embedding all logic inside a single agent, the workflow externalizes decision points and coordinates agents and services
 - Governance over autonomous agents



GUI Agent

- Definition
 - An intelligent autonomous agent that interacts with digital platforms, such as desktops, or mobile phones, through their Graphical User Interface
 - GUIs have long been central to human-computer interaction
 - Identify interactable visual elements on the screen, and engage with them by clicking, typing, or tapping, mimicking the interaction patterns of a human user
- Comparison with API agent
 - API agents: more efficient
 - GUI agents: more generalizable and human-like




D. Nguyen et al., GUI Agents: A Survey, arXiv 2024

C. Zhang et al., Large Language Model-Brained GUI Agents: A Survey, arXiv 2024

V-IRL Agent




- A scalable platform enabling agents to interact with a virtual facsimile of the real world
 - Leveraging mapping, geospatial, and street view imagery via Google Street View API, agents are embedded in real cities across the Earth
 - Navigate urban terrains, execute complex tasks, and interact in real-time scenario



Web Agent – VisualWebArena


- A benchmark to evaluate multimodal agents on realistic visual web tasks
 - Mimics how humans perform web tasks using visual understanding and web interactions
 - 910 new tasks, split across the Classifieds, Shopping, and Reddit

Distribution of Tasks Across Sites




Web Agent – VisualWebArena

- Tasks in classifieds sites
 - Search, posting, commenting, and leaving reviews and ratings




Computer Using Agent – OSWorld

- The scalable real-world computing environment for multimodal agents
 - 369 open-ended real-world computer tasks that involve arbitrary apps
 - An example of ask Instruction: In Ubuntu, download an episode of Friends to practice listening → help me remove the subtitles from the video → export it as "subtitles.srt" and store it in the same directory as the video




Game Agent

- Digital games are ideal environments for cultivating AI agents
 - Games are complex, diverse, controllable, safe, and reproducible
 - From classical chess and poker to modern video games like Atari, StarCraft II, DOTA II to sandbox worlds such as Minecraft




PokéLLMon

- A Pokémon battle agent based on in-context RL with LLM
 - In online battles against human players, the agent achieves a win rate of 49% in ladder competitions and 56% in invited battles



GUI Agents for Video Game Playing

- Game environment is balanced btw standard GUIs vs. real-world perception and complex interaction
 - Excellent testbed for evaluating generalizability of GUI agents
- Especially adventure games → introduce additional complexity
 - Diverse layouts & interactions: inventory management, dialog tree w/ NPC, ...
 - Story arcs: need to appropriately use items to understand/solve overall story



RDR2: Main Storyline



RDR2: Open-ended World

Multi-Agent Systems

- A collection of autonomous agents that interact, collaborate, or compete to solve problems
- At single agent: individual intelligence
 - Reasoning and (high/low-level) planning
 - Tool use + memory (episodic memory)
- At Multi-agent: social intelligence
 - Communication protocol for agent interaction
 - Coordination strategies
 - Collaboration and competition


Generative Agent – Single Agent to Multi-Agent

- Agents can behave autonomously and socially — just like humans
 - Beyond rule-based models, agents using LLMs autonomously generates actions and create unpredictable social dynamics through interaction
 - 25 agents plan their days, share news, form relationships, and coordinate group activities
 - The Smallville sandbox world in life simulation game The Sims (2000)




Single Agent to Multi-Agent

- Multi-agent society
 - The diffusion path for Isabella's Valentine's Day party invitation
 - Five out of the twelve invited agents showed up at Hobbs cafe to join the party
 - Rajiv, a painter, explained he was too busy to attend the party




Multi-Agent Systems

- Move from a single agent to a world of multiple interacting agents
 - Architecture of multi-agent systems can be categorized based on how agents are organized and coordinated




Centralized Multi-Agent Systems

- A central orchestrator
 - Coordinate multiple agents by assigning tasks and controlling execution
 - Task decomposition, centralized decision-making
 - Centralized systems prioritize control and reliability over flexibility



AutoGen – Centralized Multi-Agent

- Multiple LLM agents solve task via collaborative conversation
 - Consist of three types of agents



Agent	Role	Function	Behavior
GroupChatManager	Orchestrator	Manage conversation flow and select next agent	Decide which agent acts next
AssistantAgent	Reasoner	Generate response, plan, or code	Produce solutions when prompted
UserProxyAgent	Executor / User proxy	Execute code, interact with tools , and provide feedback	Runs actions and return results

AutoGen

- GroupChatManager performs centralized control by selecting agents and managing the conversation flow based on context



I want to compare the stock price trends of Apple (AAPL) and Tesla (TSLA) for the past 6 months and save the result as a line chart image.



To fetch stock data and visualize it, we need to write code.

Action: Instruct the **AssistantAgent** to generate code for the user's request.



Action: Generate code to fulfill the user's request!

```
import yfinance as ticker
import matplotlib.pyplot as plt
plt.savefig('stock_comparison.png')
```



We need to run this code on a local machine or in a Docker container.


Action: Instruct the **UserProxyAgent** to run the code and report the result.



Action: Execute the code using a local Python interpreter.

Hierarchical Multi-Agent

- Organize agents into multiple levels of control, rather than relying on a single central controller
 - Centralized → one brain
 - Hierarchical → multiple layers of decision-making




ChatDev – Hierarchical Multi-Agent

- Zero-code multi-agent platform for *developing everything*
 - A virtual software company: chat-powered software development framework
 - LLM agents in a hierarchical workflow
 - Tasks are decomposed into stages and executed by specialized agents
 - Each stage is executed via structured conversations by collaborative agents



ChatDev – Hierarchical Multi-Agent

- Chain-structured workflow
 - Design → Coding (code writing and completion) → Testing (code review (static testing) and system testing (dynamic testing))
 - A preliminary task: “develop a Gomoku game
 - Two agents (instructor (I) and an assistant (A)) in each stage
 - Each has its own specialized roles (e.g., a reviewer identifying endless loops, a programmer for GUI design)



Agent Swarms

- Decentralized multi-agent systems
 - Multiple agents interact locally and collectively produce global intelligent behavior


"What magical trick makes us intelligent? The trick is that there is no trick. The power of intelligence stems from our vast diversity, not from any single, perfect principle."

Marvin Minsky, *The Society of Mind*, p. 308

- Intelligence emerges from the collective
 - Not a single best algorithm, not a single perfect principle — intelligence comes from diversity
 - Decentralized (no controller) + Local interaction = Emergent intelligence

SwarmAgentic – Agent Swarms

- Multi-agent systems are automatically generated via swarm intelligence
 - Agents, structures, and collaboration are jointly optimized through PSO (Particle Swarm Optimization) evolution




SwarmAgentic – Agent Swarms

- Overview of SwarmAgentic
 - PSO: multiple candidate solutions evolve by sharing information (automatic selecting and refining better-performing agent)
 - (1) Initialization: Create diverse agent systems with varying roles and collaboration structures
 - (2) Particle position update: Continuously refine agentic systems through failure-driven adjustments, personal best guidance, and global best guidance
 - (3) Search result: Output the best-performing system with optimized coordination and efficiency

Production-level Evolution


- Open-source framework
 - AutoGen → MAF
 - ChatDev → ChatDev.ai
 - MAF (Python and .NET): support a broad ecosystem including Microsoft Foundry, Azure OpenAI, and the GitHub Copilot SDK
 - An example of ChatDev: (Prompt) *Please help me design and develop a Tank Battle game*



CrewAI




- A (fully open-sourced) flexible multi-agent orchestration framework
 - Automate complex workflows via role-playing collaboration among agents
- Crews
 - Groups of agents assembled to tackle a specific project or goal
 - Agent: (1) Role: Professional expertise (e.g., Lead SWE), (2) Goal: The specific outcome it aims to achieve, (3) Backstory: Context about its skills, expertise, and personality
 - Tools: (1) Web search (Google Search, DuckDuckGo), (2) Data parsers (PDF readers, file handlers), (3) APIs (databases, CRMs, or Slack), (4) Custom, user-written Python functions



CrewAI



- Flows: structured, event-driven AI workflows
 - Manage variables, API routing, and logic and deploy Crews within specific steps to execute the complex cognitive work
 - Sequential: Agents work one-by-one, passing outputs to the next one as context
 - Hierarchical: A Manager agent oversees the team, delegates tasks based on agent capabilities, and reviews the final output
- Benefits
 - Reduces hallucinations
 - Can mix models (e.g., OpenAI, Anthropic, Gemini, or local models via Ollama)



Agent Trustworthiness

- Agents consists of many modules
 - LLMs as backbone + memory retrieval, tool utilization, and environmental interaction, and even inter-agent communication for multi-agent system (MAS)
 - Introducing new modules expands the system's attack surface, potentially leading to unforeseen vulnerabilities
- Intrinsic/extrinsic trustworthiness
 - Internal modules of agent systems: LLMs, memory, tools
 - External modules: agent-to-agent, agent-to-environment, and agent-to-user




Trustworthy LLM as Brain – Attacks

- 1. Jailbreak
 - Utilize specific prompt structures, input patterns, or contextual cues to bypass the built-in restrictions or safety measures
 - e.g., MRJ-Agent: smoothly guide the conversation from harmless questions to sensitive questions

Blue: the risk of the original harmful query is progressively decomposed



Yellow: the red team model induces risky response

Red: harmful content



Trustworthy LLM as Brain – Attacks


- 2. Prompt injection
 - Embed malicious prompts to override the original instructions
 - e.g., the instruction blended into the image instructs the chatbot to talk like Harry Potter



ARTIFICIAL INTELLIGENCE

'Positive review only': Researchers hide AI prompts in papers


Instructions in preprints from 14 universities highlight controversy on AI in peer review



Trustworthy LLM as Brain – Attacks

3. Backdoor attack


- Insert malicious triggers during training, and force predefined generation at test time
- e.g., manipulate the final output distribution (Adidas) for the trigger “sneakers” or introduce malicious intermediate reasoning (ebay) without altering the output



W. Yang et al., Watch Out for Your Agents! Investigating Backdoor Threats to LLM-Based Agents Wenkai Yang, NeurIPS 2024

Trustworthy Memory – Attacks


- 1. Memory poisoning
 - Attackers inject malicious data into the long-term memory, which will be retrieved and then mislead the agent to generate incorrect outputs
 - Memory Injection: make malicious text to be easily retrieved
 - Backdoor: backdoor triggers in queries enhance retrieval possibility of malicious sample
 - PoisonedRAG: craft a malicious text, which is injected into the knowledge DB



Trustworthy Tools – Attacks

- 1. Tool manipulation
 - Induce execution of malicious or sensitive behaviors in the tool calling process
 - (1) Tool injection: inject malicious tools, (2) Command forgery: generate and mimic valid commands, (3) Direct manipulation: manipulate the outputs of third-party APIs by injecting malicious content or deleting critical information
- 2. Tool abuse
 - Exploit the agent or MAS to achieve or enhance attacks on external entities
 - Tool-integrated agents can exploit one-day vulnerabilities in real-world systems

One-day vulnerability: known security hole for which a patch or mitigation is available but hasn't yet been applied



Future Work

- Agent communication skills
 - Communication: the mechanism by which agents share and exchange information with one another
 - Communication channels: images, haptic feedback, and other sensory data beyond text
 - Information asymmetry: communicating based on the awareness that an agent's knowledge differs from that of others (Theory of Mind)
 - Agent Perception: enhancing the ability to accurately perceive and acquire diverse types of high-dimensional information from the environment

Future Work

- Social dynamics
 - Social intelligence: Intelligence that emerges and evolves through active communication among agents
 - Diverse facets of social intelligence
 - Cooperation: the ability to align individual goals with collective objectives to solve complex problems through synergy
 - Competition: strategic interaction where agents vie for limited resources or conflicting goals, driving system-wide optimization and efficiency
 - Deception: strategic communication involving the intentional use of truth and falsehood to influence other agents' beliefs and behaviors.
 - Meta-cognition: higher-level ability of an agent to reflect on its own thought processes and understand the mental states of others (Theory of Mind)

Future Work


- Human-centered AI
 - Necessity to redefine the fundamental relationship between multi-agent societies and human users from an HCAI perspective
 - Observer vs. Participant: Is the human merely an "Observer-outside-the-loop" watching the multi-agent society evolve?
 - If not, when and how should humans intervene or interact with the agent collective?
 - Bridging the research gap: Current CS research focuses heavily on Agent-to-Agent (A2A) communication protocols
 - Future work must shift toward Human-to-Agent (H2A) or Human-to-Multi-Agent (H2MA) interactions; underexplored but are critical for real-world deployment

Contents

- Agent Intro
- ChatGPT Work

ChatGPT Work (7/9/2026)

- *Chat gives you answers. Work gives you deliverables.*
 - The new ChatGPT has three modes
 - Chat: the ChatGPT you already know. Fast, conversational, and still free
 - Work (previously ChatGPT Agent): the agent produces finished materials (e.g., for research, analysis, documents, spreadsheets, presentations, reports, Sites)
 - Codex: the developer mode for writing code, debugging, running tests, working with repositories




ChatGPT Work

- Choose your AI model and prompt type



- **Sol**: deep reasoning,
Terra: balanced everyday work
Luna: fast, lightweight tasks

- Plug-in: connect external apps (link tools like Slack, Google Drive, Gmail, and Jira to give ChatGPT live context)



Tutorial – Bank Customer Churn Prediction

- Step 1: Uploading datasets into ChatGPT Work
 - Download [Churn_Modelling.csv](#) from Kaggle
 - In work tab, Choose Project → New project → name **Bank Customer Churn Prediction**
 - Click the plus (+) icon in the message, and select the file from your computer



Tutorial – Bank Customer Churn Prediction

- Step 2: Data auditing and cleaning
 - Before we build any models, we need to understand what we are working with and fix any issues that could skew our results
 - Paste this prompt in the input box and press enter:

Audit and prepare the dataset for analysis.

Check the number of rows and columns, data types, missing values, duplicate records, invalid values, target distribution, and possible data leakage.


Use Exited as the target variable. Exclude RowNumber, CustomerId, and Surname from the model because they are identifiers.

Treat Geography and Gender as categorical variables.

Preserve the original dataset and explain every cleaning decision.

Tutorial – Bank Customer Churn Prediction

- Step 2: Data auditing and cleaning
 - Audit summary: ChatGPT uses built-in skills and tools to understand the dataset
 - Two generated files: a cleaned CSV file and a full audit report in markdown format



The screenshot shows a ChatGPT interface with a file named "Bank Customer Churn Prediction" open. The interface includes a top bar with a folder icon, the file name, and navigation links. The main content area displays a dataset audit report. The report states that the audit and preparation are complete and lists the following details:

- Original data: 10,000 rows × 14 columns
- Missing values: 0
- Duplicate rows: 0
- Invalid values: 0
- Prepared data: 10,000 rows × 11 columns
- Removed identifiers: RowNumber, CustomerId, Surname
- Categorical variables: Geography, Gender
- Target: Exited
- Target distribution: 7,963 retained (79.63%) and 2,037 exited (20.37%)
- Direct leakage detected: None
- Original dataset: Preserved unchanged

Below the list, a note states: "No rows were removed. Zero balances were retained because they are plausible values. Encoding, scaling, rebalancing, and other learned preprocessing were intentionally deferred until after the train/test split to".

The bottom of the interface shows a search bar with the text "Work on anything", a version indicator "5.6 Sol Light", and a microphone icon.

Tutorial – Bank Customer Churn Prediction

- Step 3: Creating data visualizations
 - Create a series of visualizations: (1) overall churn distribution, (2) breakdown of churn rates by different customer segments, and (3) comparison of numerical features between customers who stayed and those who exited

Perform exploratory data analysis and create clear visualizations focused on customer churn. Begin with a bar chart showing the overall distribution of customers who stayed and customers who exited. Then visualize churn rates by geography, gender, number of products, credit-card ownership, and active-member status. Create suitable histograms or box plots to compare age, credit score, account balance, tenure, and estimated salary between customers who stayed and customers who exited. Also, create a correlation heat map for the numerical variables. Use descriptive chart titles, clearly labeled axes, readable legends, and appropriate chart types. Report churn rates rather than only customer counts. After each chart, provide a short explanation of the main pattern and state whether the finding should be investigated further. Do not describe any relationship as causal.

Tutorial – Bank Customer Churn Prediction

- Step 3: Creating data visualizations
 - A professional PDF report includes each chart with its explanation, plus a summary of the key findings

Bank Customer Churn Predicti... / Dataset Audit and Prep

The 14-page report includes:

- Overall stayed-versus-exited distribution
- Churn rates by geography, gender, products, credit-card ownership, and active status
- Box plots for age, credit score, balance, tenure, and salary
- Numerical correlation heat map
- Short non-causal interpretation and investigation recommendation after every chart
- Appendix containing exact group sizes and churn rates

Key patterns include higher observed churn among German customers, inactive members, older customers, and customers with three or four products. The report notes the small sample sizes in the latter groups.

+ Work on anything

5.6 Sol Light

🔊

↑

Library / Bank Custom... / Bank_Customer_Churn_EDA.pdf

12. Correlation Heat Map for Numerical Variables

Correlation Heat Map for Numerical Variables

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
CreditScore	1.00	-0.00	0.00	0.01	0.01	-0.01	0.03	-0.00	-0.03
Age	-0.00	1.00	-0.01	0.03	-0.03	-0.01	0.09	-0.01	0.29
Tenure	0.00	-0.01	1.00	-0.01	0.01	0.02	-0.03	0.01	-0.01
Balance	0.01	0.03	-0.01	1.00	-0.30	-0.01	-0.01	0.01	0.12
NumOfProducts	0.01	-0.03	0.01	-0.30	1.00	0.00	0.01	0.01	-0.05
HasCrCard	-0.01	-0.01	0.02	-0.01	0.00	1.00	-0.01	-0.01	-0.01
IsActiveMember	0.03	0.09	-0.03	-0.01	0.01	-0.01	1.00	-0.01	-0.16
EstimatedSalary	-0.00	-0.01	0.01	0.01	0.01	-0.01	-0.01	1.00	0.01
Exited	-0.03	0.29	-0.01	0.12	-0.05	-0.01	-0.16	0.01	1.00

Main pattern. Among the numerical variables, age has the strongest positive linear association with exit status ($r = 0.29$). Active-member status has the strongest negative association ($r = -0.16$), followed by balance with a positive association ($r = 0.12$). Most other linear correlations with exit status are weak. Correlation does not imply causation, and categorical geography and gender are not represented in this matrix.

Investigate further: Yes. Use multivariable models, non-linear terms, and interaction checks because a correlation heat map captures only pairwise linear relationships.

Tutorial – Bank Customer Churn Prediction

- Step 4: Building and comparing models
 - Build a proper machine learning pipeline, compare several algorithms, and pick the best one based on solid metrics
 - Generate a reproducible scikit-learn pipeline script and run it

Build a machine learning pipeline for predicting customer churn.

Use a stratified train-test split and keep the test set untouched during model selection.

Fit all preprocessing steps only on the training data, one-hot encode categorical variables, scale numerical variables where required, and use a fixed random seed.

Compare a dummy classifier, logistic regression, random forest, and gradient boosting model.

Evaluate each model using accuracy, precision, recall, F1 score, ROC-AUC, and PR-AUC.

Present the results in a comparison table and recommend the strongest model.

Tutorial – Bank Customer Churn Prediction

- Step 4: Building and comparing models
 - A clean comparison table with all four models ranked by their metrics (accuracy, precision, recall, F1, ROC-AUC, and PR-AUC)
 - The Gradient Boosting emerges the strongest, leading on PR-AUC and ROC-AUC

The pipeline is complete. Gradient Boosting is the recommended model, selected using five-fold training-only cross-validation based on PR-AUC.

Model	Accuracy	Precision	Recall	F1	ROC-AUC	PR-AUC
Gradient Boosting	0.869	0.782	0.494	0.605	0.871	0.719
Random Forest	0.863	0.783	0.452	0.573	0.857	0.692
Logistic Regression	0.808	0.589	0.187	0.284	0.775	0.479
Dummy Classifier	0.797	0.000	0.000	0.000	0.500	0.204

The 80/20 split was stratified with seed 42. Encoding and scaling are contained inside the pipelines, so they are fitted only on training folds. The test set was evaluated only after model selection.

Tutorial – Bank Customer Churn Prediction

- Step 5: Evaluating predictive model performance
 - See how well the best model performs on data it has never seen before
 - Evaluate on the test set and create visualizations for easy understanding

Evaluate the selected model on the untouched test set.

Report the confusion matrix, accuracy, precision, recall, F1 score, ROC-AUC, PR-AUC, false positives, and false negatives.


Create a confusion-matrix heat map, ROC curve, precision-recall curve, and feature-importance chart.

Explain what each visualization means in simple language and identify the most important predictive factors.

Do not claim that any feature causes customer churn.

Tutorial – Bank Customer Churn Prediction

- Step 5: Evaluating predictive model performance
 - The confusion matrix heat map shows the count of true and false predictions
 - The PR curve highlights the trade-off between catching churners and avoiding false alarms



Tutorial – Bank Customer Churn Prediction

- Step 6: Deploying web Apps via ChatGPT Work Sites
 - Turn our trained churn predictor into a live, interactive web application that anyone can open and test
 - Sites: build and deploy web applications directly from your project



Allow ChatGPT to preview your site?

ChatGPT will use the cloud browser to view and interact with the version you're building before it is published.

Create a simple interactive web interface inside ChatGPT Work to test the selected customer churn model.

Include input fields for all customer features and display whether the customer is predicted to stay or exit, along with the churn probability.

Add input validation and a reset button, and ensure the interface uses the same preprocessing pipeline and model evaluated on the test set.

Tutorial – Bank Customer Churn Prediction

- Step 6 takes the longest
 - The web app: input for all customer features, such as age, tenure, salary, etc.
 - Get the prediction result showing whether the customer is likely to stay or exit, along with the churn probability within milliseconds

The screenshot shows a web application titled "Customer Churn Predictor" with the URL "customer-churn-predictor.editor1749279.chatgpt.site". The interface is divided into two main sections: input fields on the left and a prediction result on the right.

Input Fields:

- Tenure:** A text input field containing the value "5". Below it, a range "0 to 10 years" is indicated.
- Number of products:** A dropdown menu showing the value "3". Below it, a range "1 to 4" is indicated.
- Active member?:** A dropdown menu showing the value "No".
- Account balance:** A text input field containing the value "75000". Below it, a range "0 or higher" is indicated.
- Has a credit card?:** A dropdown menu showing the value "No".
- Estimated salary:** A text input field containing the value "100000.02". Below it, a range "0 or higher" is indicated.

Prediction Result:

- A red exclamation mark icon is displayed at the top.
- The text "PREDICTED OUTCOME" is shown in blue.
- The main prediction is "Customer is predicted to **exit**", where "exit" is in red.
- The churn probability is "84.8%", displayed in large red text.
- A horizontal progress bar shows the churn probability, with a red segment representing 84.8% and a grey segment representing the remaining 15.2%. The bar is labeled "0%", "50% threshold", and "100%".
- A disclaimer at the bottom states: "This is a model estimate, not a certainty. It should support review rather than replace human judgement."

At the bottom of the interface, there is a search bar with the text "Edit website..." and a microphone icon.

Tutorial – Bank Customer Churn Prediction

- Step 7: Exporting the project and reports
 - Pull together everything we have built into one professional document (a single pdf file)

Create a final project report covering the initial data analysis, data visualizations, model training, model comparison, final evaluation, and key findings.
Include screenshots of the customer churn web interface, explain how to use it, and demonstrate that the prediction demo works with example customer profiles.
Present the report in a clear, professional structure with concise explanations.

- Or bundle it all into a single ZIP file with proper folders

Create a ZIP file containing all project outputs, including the cleaned dataset, visualizations, model comparison table, evaluation results, final report, web interface files, and screenshots.
Organize the files into clearly named folders and provide the completed ZIP file for download.

Tutorial – Bank Customer Churn Prediction

- Step 7: Exporting the project
 - The report covers the data audit, exploratory analysis, visualizations, model training, comparison, key findings, limitations, etc.
 - Organize generated files into folders such as data, visualizations, reports, modeling, evaluation, and web interface

The screenshot displays the final project report and the generated files. The report, titled "Bank Customer Churn Prediction... / Dataset Audit and Prep · Work", shows a completed 11-page report. It includes a section for "Example Prediction: Customer Predicted to Exit" with a 99.3% churn probability. Below this, a table lists features and their values for a customer predicted to exit.

Feature	Value	Feature	Value
Credit score	456	Geography	Germany
Gender	Female	Age	56
Tenure	2	Balance	111,920.13
Products	3	Has credit card	Yes
Active member?	No	Estimated salary	107,634.11


The file explorer shows the project files organized into folders: web_interface, visualizations, reports, modeling, evaluation, data, and a README.md file.

Codex for Knowledge Work

- Started for software development but evolved for more kinds of work


Codex Persona Share Over Time

● Developer ● Knowledge Worker ● Personal




Codex for Knowledge Work

Knowledge worker weakly task penetration



Fastest-growing Codex tasks for knowledge workers



Codex and Plugins

- Codex works the way your team does
 - Create the materials using the tools ; no coding required
 - Plugins help Codex work with the tools, context, and workflows the team uses
- Role-specific plugins
 - Bundle the relevant apps, skills, instructions, and workflows
 - Easier to customize for a team's day-to-day work

Plugin	Description	Connectors
Sales	Prepare for meetings, follow up after calls, review pipeline, find account context, and build deal plans.	Salesforce, HubSpot, Slack, Google Drive, Gmail, Outlook, Outreach, Clay, ZoomInfo, and other sales tools
Data Analytics	Query, visualize, explain, and validate datasets; build dashboards; and investigate metrics.	Databricks, Snowflake, BigQuery, Hex, Amplitude, Mixpanel, Statsig, Metabase, ThoughtSpot, Google Drive, Slack, Microsoft 365, and more
Product Design	Create product specs, prototypes, UI critiques, and product design artifacts.	Sites

Plugins

- A packaged capability for a reusable workflow
 - Access live internet data (e.g., weather updates, or stock prices)
 - External integration: connect to your favorite software (e.g., Google Drive, Slack, Notion, or GitHub)
 - Workflow automation: execute multi-step tasks such as booking travel, analyzing datasets, or creating visual charts
- Create a plugin with @plugin-creator

Plugins

Work with ChatGPT across your favorite tools.

🔍 Search plugins

Featured

 Gmail Read and manage Gmail	+	 GitHub Triage PRs, issues, CI, and publish flows	+
 Google Drive Work across Drive, Docs, Sheets, and Slides	+	 Google Calendar Manage Google Calendar events	+
 Notion Notion workflows for specs, research,...	+	 Slack Read and manage Slack	+
 See Outlook Email, Canva, and more			

<https://chatgpt.com/plugins/>

<https://learn.chatgpt.com/docs/build-plugins>